

On studying the performance of Hadoop Map Reduce vs MPI for Aggregation Operations: A Big Data Challenge

- Devang Swami (0625694)

I. Introduction

IT Revolution has been chiefly responsible for generating a tremendous amount of data, often referred to as Big Data. Parallel and distributed computing systems have largely been utilized for solving a problem involving big data. Of various distributed & parallel computing systems like vector processing systems, coprocessors (e.g. GPU's), and others; Map Reduce and MPI (Message Passing Interface) have found increased attention in recent times. MPI has been a de facto standard of parallel programming for decades [11], while Map Reduce is of comparatively recent development. The major difference between both the platforms is that Map Reduce uses HTTP/RPC protocols for communication which are clearly slower than the MPI's communication (which provides both point-to-point and collective communication). On the flip side, programming in MPI is clearly more difficult and requires lots of skills (in programming, networking, domain knowledge, and others) than Map Reduce. Both of these programming paradigms have been extensively utilized.

For the proposed work, minimum and maximum fare values by month & year would be calculated using Map Reduce and minimum fare by month & year would be calculated using MPI for taxi trip records data collected by New York's Taxi and Limousine Commission. The choice of maximum and minimum values by month & year was based on the fact that aggregation tasks are required by both the scientific and business community since they provide summaries of the data before proceeding with any type of analysis or visualization. The challenge here is to optimize the performance of MPI and Map Reduce for aggregation tasks. The parameters for improvement would be wall clock time and amount of data transfer in intermediate stages, as applicable.

In the following section, a review of past works has been presented. Subsequently, the problem definition has been restated. Section 4 introduces the dataset used for the project. Finally, a discussion and conclusion of the proposed work have been presented.

II. Literature Review

It is a well-established fact that Map Reduce finds application in the business community while MPI has been used more by the research community. There many available implementations of both the technologies to choose from. For instance, available MPI implementations are Open MPI, Boost MPI, MPICH, and others, while available Map Reduce implementations are Apache Hadoop, MongoDB (the database has a built-in map reduce framework.), Clusterpoint XML (another database having a built-in map reduce), and others.

The choice of implementation for MPI was Intel MPI library since it was available on SHARCNET consortium. For Map Reduce, [15] showed that Hadoop Map Reduce performs better than MongoDB build-in Map Reduce and the support for Clusterpoint seems no longer available [14]. Thus, Apache Hadoop would be the solution to test the Map Reduce framework performance which has also been a de facto tool to store big data.

III. Problem Definition

The proposed work focuses on optimizing the aggregate queries for Map Reduce Framework and MPI. Both of these models solve problems in different approaches. However, it might be noted that MPI is a paradigm and may be used to create a framework like Map reduce for instance the one by Sandia Labs [1]. There are two technologies for which optimization needs to be carried out, as follows:

- 1) Message Passing Interface: Aggregation task for the given problem may be carried out in one of the following two ways both of which may be expressed as a famous Drinking Philosophers problem [12]:

Approach 1. The Master node reads a data chunk and distributes it to the slave nodes for processing. The slave than returns the processed data. The merit of the technique is that it uses less total memory than a sequential process where master reads all the data before distributing. This technique was demonstrated in SHARCNET [13]. The following technique is more suitable in presence of a very large size of a file. Also, this program can be worked around to take advantage of replication so that processors access data file from more than one disk and decrease the I/O waiting time.

Approach 2. The master node sends a list of file names for slaves to process. Each slave processes the files provided to the same and returns the results. This approach works only if data is to be read from multiple files and if the MPI

infrastructure supports the operation of reading multiple files at the same time. The approach saves a lot of time by minimizing the time to transfer data. The only difference in regard to Approach 1, is that instead of a chunk of data we share different files. This is more suitable also in files are on separate disks. Either, it can take advantage of sharding explicitly.

Now as discussed below in Section 4, our dataset has multiple files and the size of the files are almost the same. Thus, for the given problem it is possible to use Approach 2. However, it has yet to be established if the MPI implementation at Graham or ORCA systems at SHARCNET provides the facility to access multiple files at the same time. If it is not possible to access the same file multiple times then we would shift to Approach 1.

- 2) Map Reduce: Map Reduce usually works the best when you wish to implement one-pass algorithms. For the given problem, we have to find the minimum and maximum value grouped by an attribute. Hence, no programming optimization is required for the Map Reduce code. However, there are several factors such as the amount of data chunk read for each task (AKA input split), size of the data block, and the choice of intermediate compression format; that can impact the processing time of Map Reduce tasks. Thus, we have the following three parameters for the Apache Hadoop Map Reduce implementation to optimize:

- 1) Optimal size of the block: A block (AKA physical record) is a sequence of bytes or bits, usually containing some whole number of records. In Apache Hadoop, the recommended minimum block size is 64MB and may be increased to higher size. However, since each task is carried out on a block of data, a single failure may result into redoing the work for the data of the size of one block. Hence, keeping a large block size may not be ideal. For the given problem we wish to compare the time consumption for aggregation task when block size is 64MB and 128MB.
- 2) Optimal input-split size: While processing the data in blocks, each block is split into manageable sizes called input-split. The size of input-split is a critical factor that determines the total main memory that may be utilized by the Map Reduce function. The decision regarding input-split size is based on the minimum values set in the configuration files and a heuristic that determines a good size based on block size and

other parameters or as set by the programmer/developer. These parameters may be controlled easily by using relevant functions in java Map Reduce API.

- 3) Optimal compression format: Data transmission using intermediate stages like between Mapper and Reducer tasks occur by re-writing output of the Mapper onto disk. Since disk operations are 300 times slower than the main memory operations, compressing the data during intermediate stages may save both space and time. Since data needs to be retrievable even after selecting only a portion of the compressed data. It becomes essential to compress using only splittable compressing techniques like LZO, bzip2, or snappy. Hence, for the given work we wish to test the optimal compression format for no compressing vs snappy compression technique.

The parameters for comparison for the above technologies are wall clock time for completion and amount of data written to disk during intermediate stages (only for Map Reduce).

IV. Dataset

To serve the purpose of the given work, taxi trip record dataset from New York's Taxi and Limousine Commission (TLC) for yellow taxies has been used [3]. The dataset has information about vendor id, total fare with breakups, co-ordinates for pick and drop-off, and pick up and drop of time. The data is available as files storing monthly data that begins from January, 2009 and runs till June, 2018. Total size of the dataset is 215 GB. For the given problem, we would calculate the minimum and maximum fare by month and year for the Map Reduce and maximum value for the MPI problem.

V. Discussion:

Use of distributed & parallel programming platforms have been known to require a considerable efforts in getting the right libraries, debugging network problems, safety against race conditions, and implementation of data distribution logic. For the proposed work, we have two sets of techniques wherein the said efforts would be required. Both of these techniques need a different form of skillsets and expertise. While Map Reduce requires competency in programming language java (for better control over the jobs which is available through only Java Map Reduce API), handling library issues (majorly library version problems or corrupted library builds) and YARN based scheduler policy issues

(especially in allocating virtual memory and main memory); MPI requires expertise in programming language C/C++, handling race conditions, and debugging networking skills.

For the proposed work, so far following issues have been encountered and few which have been solved:

For Map Reduce:

- 1) Oracle Java version problem: Apache Hadoop build process displayed error because Java 9 removed the activation package as a part of its main release. Initially, the problem was solved by using JAVA_OPTS as -D.javax.activation. However, later due to consideration of issues of the same package for running the software, Java 9 was degraded to Java 8.
- 2) Compression Library Issues: There were issues encountered during the build process of Snappy library support with Apache Hadoop build on Windows 10. This was solved by custom building the snappy library using Microsoft Visio. However, the build errors with snappy still prevailed. Finally, the system was shifted to Linux Mint for hosting the Apache Hadoop solution. This worked well and passed the initial debugging tests for the snappy library and has not given any warnings or errors related to Snappy library for four Map Reduce job tests. These tests were carried out using Map Reduce standard word count program.
- 3) Memory issues: Virtual memory limit issues had been encountered while performing Map Reduce on relatively small data of 700 MB. The culprit was identified to be a YARN policy. YARN does not allow virtual memory to exceed the size of main memory reserved. This was handled by suppressing the test of virtual memory calculation by setting the “yarn.nodemanager.vmem-check-enabled” flag as false. This is because it was found that enough virtual memory is available (The system has a swap space of 8GB of which less than 3GB was utilized as found using the “top” command and confirmed using “VmStat”). Another work around the same was found to be using the right value for “yarn.nodemanager.vmem-pmem-ratio”. It was identified that 2.1 usually works as the best estimate for “yarn.nodemanager.vmem-pmem-ratio” and the same may be considered in the future if a related situation arises in the future.

- 4) Storage size estimation: The amount of data to be analyzed is 215 GB for Map Reduce. However, an estimation of the size of the data during intermediate stages must be carried out to make sure that there is enough space for the intermediate data.
- 5) Data ingestion problem. Since we have data in 102 files to be loaded into HDFS before being able to perform analysis, it becomes essential to have software to ingest the data. From the literature we found multiple ingestion software's like Apache Gobblin, Apache NIFI, Apache Kafka, and others [16]. From a review of these solutions, it was found that most of them use Web HDFS to ingest data and hence it was decided that a layer of abstraction by using another software is not a good way to ingest the data. Hence, data would be ingested by using the traditional approach using the hdfs command-line tools.

For MPI:

- 1) Debugging challenges especially for race conditions and unsafe codes.
- 2) Memory issues: Usually segmentation faults and data access issues outside of the scope.
- 3) Algorithm for chunking data files: It would be essentially difficult to program an algorithm that may read a chunk of data and process it. Finally, it was found that MPI implementation at SHARCNET allows the same.
- 4) Due to limitations of the penalty imposed by SHRCNET on long-running jobs, I wish to cut down the dataset for processing of MPI to 5 GB. If time permits, analysis with even larger dataset size may be carried out.

Time Management Issues:

A lot of tasks must be accomplished in the given project and the same would be carried out using plan as shown in Figure 1.

VI. Conclusion

To tie a bow, an aggregate query to find the minimum and the maximum value of fare by month and year for Map Reduce & the minimum fare value for MPI has been proposed to study the performance & identify the best parameters for these systems. For the same after review MPICH over SHARCNET and Apache Hadoop for Map Reduce have been

suggested for measuring the performance. The parameters for comparison suggested are wall clock time for completion and amount of data written to disk during intermediate stages (only for Map Reduce). For performing the proposed analysis, NYC TLC taxi trip records dataset has been suggested, which is 215 GB in size.

Tasks\Month		October			November			
		2nd Week	3rd Week	4th Week	1st Week	2nd Week	3rd Week	4th Week
MPI SHARCNET Tasks	Optimal no of Processors		Y					
Map Reduce	Optimal block size	Y						
	Optimal input-split size		Y					
	Optimal Compression format			Y				
Compile Results					Y			
Report Writing					Y			

Figure 1. Plan for the project

Bibliography

1. Map reduce, Sandia Labs, <http://mapreduce.sandia.gov/>
2. Apache Hadoop, Hadoop Documentation, <https://hadoop.apache.org/docs>
3. NYC Taxi & Limousine Commission, New York http://www.nyc.gov/html/tlc/html/about/trip_record_data.shtml
4. Lee, K. H., Lee, Y. J., Choi, H., Chung, Y. D., & Moon, B. (2012). Parallel data processing with MapReduce: A Survey. ACM SIGMOD Record, 40(4), 11-20.
5. Dittrich, J., & Quiané-Ruiz, J. A. (2012). Efficient big data processing in Hadoop MapReduce. Proceedings of the VLDB Endowment, 5(12), 2014-2015.
6. Doukeridis, C., & Nørsvåg, K. (2014). A survey of large-scale analytical query processing in MapReduce. The VLDB Journal—The International Journal on Very Large Data Bases, 23(3), 355-380.
7. Chen, C. P., & Zhang, C. Y. (2014). Data-intensive applications, challenges, techniques and technologies: A survey on Big Data. Information Sciences, 275, 314-347.
8. Wes Kendall, Launching an Amazon EC2 MPI Cluster. <http://mpitutorial.com/tutorials/launching-an-amazon-ec2-mpi-cluster/>

9. Lu, X., Liang, F., Wang, B., Zha, L., & Xu, Z. (2014, May). DataMPI: extending MPI to hadoop-like big data computing. In Parallel and Distributed Processing Symposium, 2014 IEEE 28th International (pp. 829-838). IEEE.
10. Kang, Sol Ji, Sang Yeon Lee, and Keon Myung Lee. "Performance comparison of OpenMP, MPI, and MapReduce in practical problems." *Advances in Multimedia* 2015 (2015): 7.
11. Jin, Hui, and Xian-He Sun. "Performance comparison under failures of MPI and MapReduce: An analytical approach." *Future Generation Computer Systems* 29.7 (2013): 1808-1815.
12. Kshemkalyani, A. D., & Singhal, M. (2011). *Distributed computing: principles, algorithms, and systems*. Cambridge University Press.
13. Sharcnet HPC. (Sept, 18, 2018). Concurrent File I/O by Multiple Processes. <https://www.youtube.com/watch?v=3s-dBEopfWQ>
14. Dbengines. Clusterpoint System Properties. <https://db-engines.com/en/system/Clusterpoint>
15. Dede, E., Govindaraju, M., Gunter, D., Canon, R. S., & Ramakrishnan, L. (2013, June). Performance evaluation of a mongodb and hadoop platform for scientific data analysis. In *Proceedings of the 4th ACM workshop on Scientific cloud computing* (pp. 13-20). ACM.
16. PAT Research. Top 18 Data ingestion tools. <https://www.predictiveanalyticstoday.com/data-ingestion-tools/>